

aai-520-final-project-workflow

October 23, 2023

0.0.1 Install Python Libraries

```
[55]: !pip install kaggle
      !pip install tensorflow
      !pip install transformers
      !pip install evaluate
      !pip install scikit-learn
      !pip install datasets
      !pip install seaborn
      !pip install matplotlib
      !pip install typing
      !pip install pandas
```

```
Requirement already satisfied: kaggle in
/Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-
packages (1.5.16)
Requirement already satisfied: six>=1.10 in
/Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-
packages (from kaggle) (1.16.0)
Requirement already satisfied: certifi in
/Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-
packages (from kaggle) (2023.7.22)
Requirement already satisfied: python-dateutil in
/Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-
packages (from kaggle) (2.8.2)
Requirement already satisfied: requests in
/Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-
packages (from kaggle) (2.31.0)
Requirement already satisfied: tqdm in
/Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-
packages (from kaggle) (4.65.0)
Requirement already satisfied: python-slugify in
/Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-
packages (from kaggle) (8.0.1)
Requirement already satisfied: urllib3 in
/Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-
packages (from kaggle) (1.26.16)
Requirement already satisfied: bleach in
```

/Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (from kaggle) (6.1.0)
Requirement already satisfied: webencodings in
/Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (from bleach->kaggle) (0.5.1)
Requirement already satisfied: text-unidecode>=1.3 in
/Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (from python-slugify->kaggle) (1.3)
Requirement already satisfied: charset-normalizer<4,>=2 in
/Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (from requests->kaggle) (2.0.4)
Requirement already satisfied: idna<4,>=2.5 in
/Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (from requests->kaggle) (3.4)
Requirement already satisfied: tensorflow in
/Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (2.12.0)
Requirement already satisfied: absl-py>=1.0.0 in
/Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (from tensorflow) (1.4.0)
Requirement already satisfied: astunparse>=1.6.0 in
/Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (from tensorflow) (1.6.3)
Requirement already satisfied: flatbuffers>=2.0 in
/Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (from tensorflow) (2.0)
Requirement already satisfied: gast<=0.4.0,>=0.2.1 in
/Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (from tensorflow) (0.4.0)
Requirement already satisfied: google-pasta>=0.1.1 in
/Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (from tensorflow) (0.2.0)
Requirement already satisfied: h5py>=2.9.0 in
/Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (from tensorflow) (3.7.0)
Requirement already satisfied: jax>=0.3.15 in
/Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (from tensorflow) (0.4.18)
Requirement already satisfied: libclang>=13.0.0 in
/Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (from tensorflow) (16.0.6)
Requirement already satisfied: numpy<1.24,>=1.22 in
/Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (from tensorflow) (1.23.5)
Requirement already satisfied: opt-einsum>=2.3.2 in
/Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (from tensorflow) (3.3.0)
Requirement already satisfied: packaging in

```

/Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-
packages (from tensorflow) (23.0)
Requirement already satisfied:
protobuf!=4.21.0,!4.21.1,!4.21.2,!4.21.3,!4.21.4,!4.21.5,<5.0.0dev,>=3.20.3
in /Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-
packages (from tensorflow) (3.20.3)
Requirement already satisfied: setuptools in
/Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-
packages (from tensorflow) (68.0.0)
Requirement already satisfied: six>=1.12.0 in
/Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-
packages (from tensorflow) (1.16.0)
Requirement already satisfied: termcolor>=1.1.0 in
/Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-
packages (from tensorflow) (2.1.0)
Requirement already satisfied: typing-extensions>=3.6.6 in
/Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-
packages (from tensorflow) (4.7.1)
Requirement already satisfied: wrapt<1.15,>=1.11.0 in
/Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-
packages (from tensorflow) (1.14.1)
Requirement already satisfied: grpcio<2.0,>=1.24.3 in
/Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-
packages (from tensorflow) (1.48.2)
Requirement already satisfied: tensorboard<2.13,>=2.12 in
/Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-
packages (from tensorflow) (2.12.1)
Requirement already satisfied: tensorflow-estimator<2.13,>=2.12.0 in
/Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-
packages (from tensorflow) (2.12.0)
Requirement already satisfied: keras<2.13,>=2.12.0 in
/Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-
packages (from tensorflow) (2.12.0)
Requirement already satisfied: wheel<1.0,>=0.23.0 in
/Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-
packages (from astunparse>=1.6.0->tensorflow) (0.35.1)
Requirement already satisfied: ml-dtypes>=0.2.0 in
/Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-
packages (from jax>=0.3.15->tensorflow) (0.3.1)
Requirement already satisfied: scipy>=1.7 in
/Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-
packages (from jax>=0.3.15->tensorflow) (1.10.1)
Requirement already satisfied: google-auth<3,>=1.6.3 in
/Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-
packages (from tensorboard<2.13,>=2.12->tensorflow) (2.6.0)
Requirement already satisfied: google-auth-oauthlib<1.1,>=0.5 in
/Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-
packages (from tensorboard<2.13,>=2.12->tensorflow) (0.5.2)

```

Requirement already satisfied: markdown>=2.6.8 in
/Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-
packages (from tensorboard<2.13,>=2.12->tensorflow) (3.4.1)

Requirement already satisfied: requests<3,>=2.21.0 in
/Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-
packages (from tensorboard<2.13,>=2.12->tensorflow) (2.31.0)

Requirement already satisfied: tensorboard-data-server<0.8.0,>=0.7.0 in
/Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-
packages (from tensorboard<2.13,>=2.12->tensorflow) (0.7.0)

Requirement already satisfied: tensorboard-plugin-wit>=1.6.0 in
/Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-
packages (from tensorboard<2.13,>=2.12->tensorflow) (1.8.1)

Requirement already satisfied: werkzeug>=1.0.1 in
/Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-
packages (from tensorboard<2.13,>=2.12->tensorflow) (2.2.3)

Requirement already satisfied: cachetools<6.0,>=2.0.0 in
/Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-
packages (from google-auth<3,>=1.6.3->tensorboard<2.13,>=2.12->tensorflow)
(4.2.2)

Requirement already satisfied: pyasn1-modules>=0.2.1 in
/Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-
packages (from google-auth<3,>=1.6.3->tensorboard<2.13,>=2.12->tensorflow)
(0.2.8)

Requirement already satisfied: rsa<5,>=3.1.4 in
/Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-
packages (from google-auth<3,>=1.6.3->tensorboard<2.13,>=2.12->tensorflow)
(4.7.2)

Requirement already satisfied: requests-oauthlib>=0.7.0 in
/Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-
packages (from google-auth-
oauthlib<1.1,>=0.5->tensorboard<2.13,>=2.12->tensorflow) (1.3.0)

Requirement already satisfied: charset-normalizer<4,>=2 in
/Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-
packages (from requests<3,>=2.21.0->tensorboard<2.13,>=2.12->tensorflow) (2.0.4)

Requirement already satisfied: idna<4,>=2.5 in
/Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-
packages (from requests<3,>=2.21.0->tensorboard<2.13,>=2.12->tensorflow) (3.4)

Requirement already satisfied: urllib3<3,>=1.21.1 in
/Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-
packages (from requests<3,>=2.21.0->tensorboard<2.13,>=2.12->tensorflow)
(1.26.16)

Requirement already satisfied: certifi>=2017.4.17 in
/Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-
packages (from requests<3,>=2.21.0->tensorboard<2.13,>=2.12->tensorflow)
(2023.7.22)

Requirement already satisfied: MarkupSafe>=2.1.1 in
/Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-
packages (from werkzeug>=1.0.1->tensorboard<2.13,>=2.12->tensorflow) (2.1.1)

Requirement already satisfied: pyasn1<0.5.0,>=0.4.6 in
/Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-
packages (from pyasn1-modules>=0.2.1->google-
auth<3,>=1.6.3->tensorboard<2.13,>=2.12->tensorflow) (0.4.8)

Requirement already satisfied: oauthlib>=3.0.0 in
/Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-
packages (from requests-oauthlib>=0.7.0->google-auth-
oauthlib<1.1,>=0.5->tensorboard<2.13,>=2.12->tensorflow) (3.2.2)

Requirement already satisfied: transformers in
/Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-
packages (4.32.1)

Requirement already satisfied: filelock in
/Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-
packages (from transformers) (3.9.0)

Requirement already satisfied: huggingface-hub<1.0,>=0.15.1 in
/Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-
packages (from transformers) (0.15.1)

Requirement already satisfied: numpy>=1.17 in
/Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-
packages (from transformers) (1.23.5)

Requirement already satisfied: packaging>=20.0 in
/Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-
packages (from transformers) (23.0)

Requirement already satisfied: pyyaml>=5.1 in
/Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-
packages (from transformers) (6.0)

Requirement already satisfied: regex!=2019.12.17 in
/Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-
packages (from transformers) (2022.7.9)

Requirement already satisfied: requests in
/Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-
packages (from transformers) (2.31.0)

Requirement already satisfied: tokenizers!=0.11.3,<0.14,>=0.11.1 in
/Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-
packages (from transformers) (0.13.2)

Requirement already satisfied: safetensors>=0.3.1 in
/Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-
packages (from transformers) (0.3.2)

Requirement already satisfied: tqdm>=4.27 in
/Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-
packages (from transformers) (4.65.0)

Requirement already satisfied: fsspec in
/Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-
packages (from huggingface-hub<1.0,>=0.15.1->transformers) (2023.9.2)

Requirement already satisfied: typing-extensions>=3.7.4.3 in
/Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-
packages (from huggingface-hub<1.0,>=0.15.1->transformers) (4.7.1)

Requirement already satisfied: charset-normalizer<4,>=2 in

/Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (from requests->transformers) (2.0.4)
Requirement already satisfied: idna<4,>=2.5 in
/Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (from requests->transformers) (3.4)
Requirement already satisfied: urllib3<3,>=1.21.1 in
/Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (from requests->transformers) (1.26.16)
Requirement already satisfied: certifi>=2017.4.17 in
/Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (from requests->transformers) (2023.7.22)
Requirement already satisfied: evaluate in
/Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (0.4.0)
Requirement already satisfied: datasets>=2.0.0 in
/Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (from evaluate) (2.12.0)
Requirement already satisfied: numpy>=1.17 in
/Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (from evaluate) (1.23.5)
Requirement already satisfied: dill in
/Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (from evaluate) (0.3.6)
Requirement already satisfied: pandas in
/Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (from evaluate) (1.5.3)
Requirement already satisfied: requests>=2.19.0 in
/Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (from evaluate) (2.31.0)
Requirement already satisfied: tqdm>=4.62.1 in
/Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (from evaluate) (4.65.0)
Requirement already satisfied: xxhash in
/Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (from evaluate) (2.0.2)
Requirement already satisfied: multiprocessing in
/Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (from evaluate) (0.70.14)
Requirement already satisfied: fsspec[http]>=2021.05.0 in
/Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (from evaluate) (2023.9.2)
Requirement already satisfied: huggingface-hub>=0.7.0 in
/Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (from evaluate) (0.15.1)
Requirement already satisfied: packaging in
/Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (from evaluate) (23.0)
Requirement already satisfied: responses<0.19 in

/Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (from evaluate) (0.13.3)
 Requirement already satisfied: pyarrow>=8.0.0 in
 /Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (from datasets>=2.0.0->evaluate) (11.0.0)
 Requirement already satisfied: aiohttp in
 /Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (from datasets>=2.0.0->evaluate) (3.8.3)
 Requirement already satisfied: pyyaml>=5.1 in
 /Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (from datasets>=2.0.0->evaluate) (6.0)
 Requirement already satisfied: filelock in
 /Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (from huggingface-hub>=0.7.0->evaluate) (3.9.0)
 Requirement already satisfied: typing-extensions>=3.7.4.3 in
 /Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (from huggingface-hub>=0.7.0->evaluate) (4.7.1)
 Requirement already satisfied: charset-normalizer<4,>=2 in
 /Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (from requests>=2.19.0->evaluate) (2.0.4)
 Requirement already satisfied: idna<4,>=2.5 in
 /Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (from requests>=2.19.0->evaluate) (3.4)
 Requirement already satisfied: urllib3<3,>=1.21.1 in
 /Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (from requests>=2.19.0->evaluate) (1.26.16)
 Requirement already satisfied: certifi>=2017.4.17 in
 /Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (from requests>=2.19.0->evaluate) (2023.7.22)
 Requirement already satisfied: six in
 /Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (from responses<0.19->evaluate) (1.16.0)
 Requirement already satisfied: python-dateutil>=2.8.1 in
 /Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (from pandas->evaluate) (2.8.2)
 Requirement already satisfied: pytz>=2020.1 in
 /Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (from pandas->evaluate) (2022.7)
 Requirement already satisfied: attrs>=17.3.0 in
 /Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (from aiohttp->datasets>=2.0.0->evaluate) (22.1.0)
 Requirement already satisfied: multidict<7.0,>=4.5 in
 /Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (from aiohttp->datasets>=2.0.0->evaluate) (6.0.2)
 Requirement already satisfied: async-timeout<5.0,>=4.0.0a3 in
 /Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (from aiohttp->datasets>=2.0.0->evaluate) (4.0.2)
 Requirement already satisfied: yarl<2.0,>=1.0 in

/Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (from aiohttp->datasets>=2.0.0->evaluate) (1.8.1)
 Requirement already satisfied: frozenlist>=1.1.1 in /Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (from aiohttp->datasets>=2.0.0->evaluate) (1.3.3)
 Requirement already satisfied: aiosignal>=1.1.2 in /Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (from aiohttp->datasets>=2.0.0->evaluate) (1.2.0)
 Requirement already satisfied: scikit-learn in /Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (1.2.2)
 Requirement already satisfied: numpy>=1.17.3 in /Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (from scikit-learn) (1.23.5)
 Requirement already satisfied: scipy>=1.3.2 in /Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (from scikit-learn) (1.10.1)
 Requirement already satisfied: joblib>=1.1.1 in /Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (from scikit-learn) (1.2.0)
 Requirement already satisfied: threadpoolctl>=2.0.0 in /Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (from scikit-learn) (2.2.0)
 Requirement already satisfied: datasets in /Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (2.12.0)
 Requirement already satisfied: numpy>=1.17 in /Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (from datasets) (1.23.5)
 Requirement already satisfied: pyarrow>=8.0.0 in /Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (from datasets) (11.0.0)
 Requirement already satisfied: dill<0.3.7,>=0.3.0 in /Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (from datasets) (0.3.6)
 Requirement already satisfied: pandas in /Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (from datasets) (1.5.3)
 Requirement already satisfied: requests>=2.19.0 in /Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (from datasets) (2.31.0)
 Requirement already satisfied: tqdm>=4.62.1 in /Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (from datasets) (4.65.0)
 Requirement already satisfied: xxhash in /Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (from datasets) (2.0.2)
 Requirement already satisfied: multiprocessing in

/Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (from datasets) (0.70.14)
 Requirement already satisfied: fsspec[http]>=2021.11.1 in /Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (from datasets) (2023.9.2)
 Requirement already satisfied: aiohttp in /Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (from datasets) (3.8.3)
 Requirement already satisfied: huggingface-hub<1.0.0,>=0.11.0 in /Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (from datasets) (0.15.1)
 Requirement already satisfied: packaging in /Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (from datasets) (23.0)
 Requirement already satisfied: responses<0.19 in /Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (from datasets) (0.13.3)
 Requirement already satisfied: pyyaml>=5.1 in /Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (from datasets) (6.0)
 Requirement already satisfied: attrs>=17.3.0 in /Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (from aiohttp->datasets) (22.1.0)
 Requirement already satisfied: charset-normalizer<3.0,>=2.0 in /Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (from aiohttp->datasets) (2.0.4)
 Requirement already satisfied: multidict<7.0,>=4.5 in /Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (from aiohttp->datasets) (6.0.2)
 Requirement already satisfied: async-timeout<5.0,>=4.0.0a3 in /Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (from aiohttp->datasets) (4.0.2)
 Requirement already satisfied: yarl<2.0,>=1.0 in /Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (from aiohttp->datasets) (1.8.1)
 Requirement already satisfied: frozenlist>=1.1.1 in /Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (from aiohttp->datasets) (1.3.3)
 Requirement already satisfied: aiosignal>=1.1.2 in /Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (from aiohttp->datasets) (1.2.0)
 Requirement already satisfied: filelock in /Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (from huggingface-hub<1.0.0,>=0.11.0->datasets) (3.9.0)
 Requirement already satisfied: typing-extensions>=3.7.4.3 in /Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (from huggingface-hub<1.0.0,>=0.11.0->datasets) (4.7.1)
 Requirement already satisfied: idna<4,>=2.5 in

/Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (from requests>=2.19.0->datasets) (3.4)
 Requirement already satisfied: urllib3<3,>=1.21.1 in
 /Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (from requests>=2.19.0->datasets) (1.26.16)
 Requirement already satisfied: certifi>=2017.4.17 in
 /Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (from requests>=2.19.0->datasets) (2023.7.22)
 Requirement already satisfied: six in
 /Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (from responses<0.19->datasets) (1.16.0)
 Requirement already satisfied: python-dateutil>=2.8.1 in
 /Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (from pandas->datasets) (2.8.2)
 Requirement already satisfied: pytz>=2020.1 in
 /Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (from pandas->datasets) (2022.7)
 Requirement already satisfied: seaborn in
 /Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (0.12.2)
 Requirement already satisfied: numpy!=1.24.0,>=1.17 in
 /Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (from seaborn) (1.23.5)
 Requirement already satisfied: pandas>=0.25 in
 /Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (from seaborn) (1.5.3)
 Requirement already satisfied: matplotlib!=3.6.1,>=3.1 in
 /Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (from seaborn) (3.7.1)
 Requirement already satisfied: contourpy>=1.0.1 in
 /Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (from matplotlib!=3.6.1,>=3.1->seaborn) (1.0.5)
 Requirement already satisfied: cycler>=0.10 in
 /Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (from matplotlib!=3.6.1,>=3.1->seaborn) (0.11.0)
 Requirement already satisfied: fonttools>=4.22.0 in
 /Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (from matplotlib!=3.6.1,>=3.1->seaborn) (4.25.0)
 Requirement already satisfied: kiwisolver>=1.0.1 in
 /Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (from matplotlib!=3.6.1,>=3.1->seaborn) (1.4.4)
 Requirement already satisfied: packaging>=20.0 in
 /Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (from matplotlib!=3.6.1,>=3.1->seaborn) (23.0)
 Requirement already satisfied: pillow>=6.2.0 in
 /Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (from matplotlib!=3.6.1,>=3.1->seaborn) (9.4.0)
 Requirement already satisfied: pyparsing>=2.3.1 in

/Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (from matplotlib!=3.6.1,>=3.1->seaborn) (3.0.9)
 Requirement already satisfied: python-dateutil>=2.7 in /Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (from matplotlib!=3.6.1,>=3.1->seaborn) (2.8.2)
 Requirement already satisfied: pytz>=2020.1 in /Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (from pandas>=0.25->seaborn) (2022.7)
 Requirement already satisfied: six>=1.5 in /Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (from python-dateutil>=2.7->matplotlib!=3.6.1,>=3.1->seaborn) (1.16.0)
 Requirement already satisfied: matplotlib in /Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (3.7.1)
 Requirement already satisfied: contourpy>=1.0.1 in /Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (from matplotlib) (1.0.5)
 Requirement already satisfied: cycler>=0.10 in /Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (from matplotlib) (0.11.0)
 Requirement already satisfied: fonttools>=4.22.0 in /Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (from matplotlib) (4.25.0)
 Requirement already satisfied: kiwisolver>=1.0.1 in /Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (from matplotlib) (1.4.4)
 Requirement already satisfied: numpy>=1.20 in /Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (from matplotlib) (1.23.5)
 Requirement already satisfied: packaging>=20.0 in /Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (from matplotlib) (23.0)
 Requirement already satisfied: pillow>=6.2.0 in /Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (from matplotlib) (9.4.0)
 Requirement already satisfied: pyparsing>=2.3.1 in /Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (from matplotlib) (3.0.9)
 Requirement already satisfied: python-dateutil>=2.7 in /Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (from matplotlib) (2.8.2)
 Requirement already satisfied: six>=1.5 in /Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (from python-dateutil>=2.7->matplotlib) (1.16.0)
 Requirement already satisfied: typing in /Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-packages (3.7.4.3)
 Requirement already satisfied: pandas in

```
/Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-  
packages (1.5.3)  
Requirement already satisfied: python-dateutil>=2.8.1 in  
/Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-  
packages (from pandas) (2.8.2)  
Requirement already satisfied: pytz>=2020.1 in  
/Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-  
packages (from pandas) (2022.7)  
Requirement already satisfied: numpy>=1.21.0 in  
/Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-  
packages (from pandas) (1.23.5)  
Requirement already satisfied: six>=1.5 in  
/Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-  
packages (from python-dateutil>=2.8.1->pandas) (1.16.0)
```

0.0.2 Apply Kaggle Configuration

0.0.3 Import Libraries

```
[58]: import numpy as np # linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)
import os
import math

import kaggle
import tensorflow as tf
from datasets import Dataset, concatenate_datasets
from typing import Union, Optional

from transformers import AutoTokenizer, AdamWeightDecay, BlenderbotTokenizer,
    ↳TFBlenderbotForConditionalGeneration,
    ↳DataCollatorForLanguageModeling,DataCollatorForSeq2Seq,DataCollatorWithPadding,
    ↳create_optimizer, TFAutoModelForCausalLM,TFAutoModelForSeq2SeqLM, pipeline
# Used for model metric call when applying the metrics during training and
    ↳evaluation
from transformers.keras_callbacks import KerasMetricCallback
from tensorflow.keras.optimizers import Adam, Adagrad, SGD

#from dataclasses import dataclass
from transformers.tokenization_utils_base import PreTrainedTokenizerBase,
    ↳PaddingStrategy

import evaluate

# Import sklearn metrics
from sklearn.metrics import accuracy_score, precision_score, recall_score,
    ↳f1_score, confusion_matrix
```

0.0.4 Import Dataset

- huggingface ubuntu dialog corpus
- https://huggingface.co/datasets/ubuntu_dialogs_corpus
- authenticate with Kaggle
- download the files from Kaggle
- add them to dataset module

```
[59]: kaggle.api.authenticate()
```

Warning: Your Kaggle API key is readable by other users on this system! To fix this, you can run 'chmod 600 /Users/hudsonshimanyula/.kaggle/kaggle.json'

```
[61]: # Define the directory where the files are located
directory = "./Ubuntu-dialogue-corpus"
```

```
[62]: # Define the file names and corresponding dataframe names
files_dataframes = {
    "dialogueText_301.csv": "train_ubuntu_dialogue_df",
    "dialogueText_196.csv": "validation_ubuntu_dialogue_df",
    "dialogueText.csv": "test_ubuntu_dialogue_df"
}
```

```
[63]: # Iterate through the files in the directory
for filename in os.listdir(directory):
    # If the file is one of the files we're interested in
    if filename in files_dataframes:
        # Construct the full file path
        file_path = os.path.join(directory, filename)
        # Load the file into a dataframe and assign it to a variable with the
        ↪corresponding name
        globals()[files_dataframes[filename]] = pd.read_csv(file_path)
```

```
[64]: train_ubuntu_dialogue_df.head()
```

```
[64]:
```

	folder	dialogueID	date	from	to	\
0	301	1.tsv	2004-11-23T11:49:00.000Z	stuNNed	NaN	
1	301	1.tsv	2004-11-23T11:49:00.000Z	crimsun	stuNNed	
2	301	1.tsv	2004-11-23T11:49:00.000Z	stuNNed	crimsun	
3	301	1.tsv	2004-11-23T11:49:00.000Z	crimsun	stuNNed	
4	301	1.tsv	2004-11-23T11:50:00.000Z	stuNNed	crimsun	

	text
0	any ideas why java plugin takes so long to load?
1	java 1.4?
2	yes
3	java 1.5 loads _much_ faster
4	noneus: how can i get 1.5 is there a .deb some...

```
[65]: validation_ubuntu_dialogue_df.head()
```

```
[65]:
```

	folder	dialogueID	date	from	to	\
0	301	1.tsv	2004-11-23T11:49:00.000Z	stuNNed	NaN	
1	301	1.tsv	2004-11-23T11:49:00.000Z	crimsun	stuNNed	
2	301	1.tsv	2004-11-23T11:49:00.000Z	stuNNed	crimsun	
3	301	1.tsv	2004-11-23T11:49:00.000Z	crimsun	stuNNed	
4	301	1.tsv	2004-11-23T11:50:00.000Z	stuNNed	crimsun	

	text
0	any ideas why java plugin takes so long to load?
1	java 1.4?
2	yes
3	java 1.5 loads _much_ faster

4 noneus: how can i get 1.5 is there a .deb some...

```
[66]: # Review the downloaded dataset
test_ubuntu_dialogue_df.head()
```

```
[66]:
```

	folder	dialogueID	date	from	to	\
0	3	126125.tsv	2008-04-23T14:55:00.000Z	bad_image	NaN	
1	3	126125.tsv	2008-04-23T14:56:00.000Z	bad_image	NaN	
2	3	126125.tsv	2008-04-23T14:57:00.000Z	lordleemo	bad_image	
3	3	64545.tsv	2009-08-01T06:22:00.000Z	mechtech	NaN	
4	3	64545.tsv	2009-08-01T06:22:00.000Z	mechtech	NaN	

```
text
0 Hello folks, please help me a bit with the fol...
1 Did I choose a bad channel? I ask because you ...
2 the second sentence is better english and we...
3 Sock Puppe?t
4 WTF?
```

0.0.5 Exploratory Data Analysis

- get dataset info
- get dataset description

```
[67]: train_ubuntu_dialogue_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 16587830 entries, 0 to 16587829
Data columns (total 6 columns):
#   Column      Dtype
---  ---
0   folder      int64
1   dialogueID  object
2   date        object
3   from        object
4   to          object
5   text        object
dtypes: int64(1), object(5)
memory usage: 759.3+ MB
```

```
[68]: validation_ubuntu_dialogue_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 9212877 entries, 0 to 9212876
Data columns (total 6 columns):
#   Column      Dtype
---  ---
0   folder      int64
```

```
1  dialogueID  object
2  date       object
3  from       object
4  to         object
5  text       object
dtypes: int64(1), object(5)
memory usage: 421.7+ MB
```

```
[69]: test_ubuntu_dialogue_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1038324 entries, 0 to 1038323
Data columns (total 6 columns):
#   Column      Non-Null Count  Dtype
---  -
0   folder      1038324 non-null  int64
1   dialogueID  1038324 non-null  object
2   date        1038324 non-null  object
3   from        1038311 non-null  object
4   to          566035 non-null  object
5   text        1038237 non-null  object
dtypes: int64(1), object(5)
memory usage: 47.5+ MB
```

```
[70]: train_ubuntu_dialogue_df['text'].describe()
```

```
[70]: count      16586581
unique      12520971
top         yes
freq        83685
Name: text, dtype: object
```

```
[71]: validation_ubuntu_dialogue_df['text'].describe()
```

```
[71]: count      9212063
unique      7232452
top         yes
freq        47000
Name: text, dtype: object
```

```
[72]: test_ubuntu_dialogue_df['text'].describe()
```

```
[72]: count      1038237
unique      863907
top         thanks
freq        6256
Name: text, dtype: object
```


0.0.6 Preprocess the datasets

- convert the dataframes to datasets
- initialize the tokenizer by selecting a pretrained model
- initialize the model by selecting a pretrained model
- initialize the datacollator
- apply tokenizer to datacollator
- tokenize the dataset with the tokenizer function
- build a tokenized dataset
- build function to split concatenated sequences into shorter chunks defined by block size and use that to build the label

```
[73]: train_ubuntu_dialogue_ds = Dataset.from_pandas(train_ubuntu_dialogue_df[0:10])
validation_ubuntu_dialogue_ds = Dataset.
↳from_pandas(validation_ubuntu_dialogue_df[0:10])
test_ubuntu_dialogue_ds = Dataset.from_pandas(test_ubuntu_dialogue_df[0:10])
```

```
[74]: train_ubuntu_dialogue_ds
```

```
[74]: Dataset({
    features: ['folder', 'dialogueID', 'date', 'from', 'to', 'text'],
    num_rows: 10
})
```

```
[75]: tokenizer = BlenderbotTokenizer.from_pretrained('facebook/
↳blenderbot-400M-distill') #AutoTokenizer.from_pretrained('facebook/
↳blenderbot-400M-distill') #, truncation=True, padding="longest",
↳max_length=512, return_overflowing_tokens=True, return_offsets_mapping=True)
↳#, truncation=True
```

```
Downloading (...)olve/main/vocab.json: 100%|      | 127k/127k [00:00<00:00,
3.60MB/s]
Downloading (...)olve/main/merges.txt: 100%|      | 62.9k/62.9k [00:00<00:00,
75.5MB/s]
Downloading (...)okenizer_config.json: 100%|      | 1.15k/1.15k [00:00<00:00,
865kB/s]
Downloading (...)in/added_tokens.json: 100%|      | 16.0/16.0 [00:00<00:00,
57.1kB/s]
Downloading (...)cial_tokens_map.json: 100%|      | 772/772 [00:00<00:00,
5.11MB/s]
Downloading (...)lve/main/config.json: 100%|      | 1.57k/1.57k [00:00<00:00,
4.03MB/s]
```

```
[76]: tokenizer.pad_token = tokenizer.eos_token
```

```
[77]: # used for causal inference
data_collator = DataCollatorForSeq2Seq(tokenizer=tokenizer, return_tensors='tf')
```

```
[78]: tokenizer
```

```
[78]: BlenderbotTokenizer(name_or_path='facebook/blenderbot-400M-distill',
vocab_size=8008, model_max_length=128, is_fast=False, padding_side='right',
truncation_side='right', special_tokens={'bos_token': AddedToken("<s>",
rstrip=False, lstrip=False, single_word=False, normalized=True), 'eos_token':
AddedToken("</s>", rstrip=False, lstrip=False, single_word=False,
normalized=True), 'unk_token': AddedToken("<unk>", rstrip=False, lstrip=False,
single_word=False, normalized=True), 'sep_token': AddedToken("</s>",
rstrip=False, lstrip=False, single_word=False, normalized=True), 'pad_token':
'</s>', 'cls_token': AddedToken("<s>", rstrip=False, lstrip=False,
single_word=False, normalized=True), 'mask_token': AddedToken("<mask>",
rstrip=False, lstrip=True, single_word=False, normalized=True)}},
clean_up_tokenization_spaces=True)
```

```
[79]: def preprocess_function(examples):
    """
    This function preprocesses the input examples for the model training. It
    ↪tokenizes the input text, removes examples where input_ids is None, and
    ↪creates decoder_input_ids and labels.

    Args:
        examples (dict): A dictionary containing the input text.

    Returns:
        dict: A dictionary containing the tokenized input text, decoder_input_ids,
    ↪and labels.
    """
    # Tokenize the examples
    tokenized_examples = tokenizer([" ".join(x) for x in examples['text']],
    ↪truncation=True, padding="longest", max_length=512)

    # Remove examples where input_ids is None
    tokenized_examples = {k: v for k, v in tokenized_examples.items() if v is
    ↪not None}

    # Create decoder_input_ids and labels
    tokenized_examples["decoder_input_ids"] = [[tokenizer.bos_token_id] + ids +
    ↪[tokenizer.eos_token_id] for ids in tokenized_examples["input_ids"]]
    tokenized_examples["labels"] = [[tokenizer.bos_token_id] + ids + [tokenizer.
    ↪eos_token_id] for ids in tokenized_examples["input_ids"]]

    return tokenized_examples
```

```
[80]: tokenized_train_ubuntu_dialogue_ds = train_ubuntu_dialogue_ds.map(
    preprocess_function,
    batched=True,
```

```

num_proc=1,
remove_columns=train_ubuntu_dialogue_ds.column_names,
)

```

Map: 0%| | 0/10 [00:00<?, ? examples/s]

```

[81]: tokenized_validation_ubuntu_dialogue_ds = validation_ubuntu_dialogue_ds.map(
    preprocess_function,
    batched=True,
    num_proc=1,
    remove_columns=validation_ubuntu_dialogue_ds.column_names,
)

```

```

[82]: tokenized_test_ubuntu_dialogue_ds = test_ubuntu_dialogue_ds.map(
    preprocess_function,
    batched=True,
    num_proc=1,
    remove_columns=test_ubuntu_dialogue_ds.column_names,
)

```

Map: 0%| | 0/10 [00:00<?, ? examples/s]

```

[83]: tokenized_train_ubuntu_dialogue_ds

```

```

[83]: Dataset({
    features: ['input_ids', 'attention_mask', 'decoder_input_ids', 'labels'],
    num_rows: 10
})

```

0.0.7 Model Build

- build the hyperparameters
- build and compute metrics for model evaluation
- build the pretrained model
- build the tensorflow prepared dataset (includes tensorflow tensor conversion, batch setting for model runs, shuffling the data, adding the data collator for collation)
- build metric callbacks
- apply them to the model
- compile the model

```

[85]: model = TFBleNDERbotForConditionalGeneration.from_pretrained('facebook/
    ↪blenderbot-400M-distill')

```

Downloading tf_model.h5: 100%| | 1.46G/1.46G [00:28<00:00, 51.8MB/s]

All model checkpoint layers were used when initializing

TFBlenderbotForConditionalGeneration.

Some layers of TFBlenderbotForConditionalGeneration were not initialized from the model checkpoint at facebook/blenderbot-400M-distill and are newly initialized: ['final_logits_bias']
You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.
Downloading (...)neration_config.json: 100%| | 347/347 [00:00<00:00, 4.58MB/s]

```
[86]: # Preparing TensorFlow dataset for training:
# The `prepare_tf_dataset` method is called on the model object to convert the
# ↪tokenized training dataset into a TensorFlow dataset.
# The dataset is shuffled to ensure randomness, which is beneficial for
# ↪training the model.
# Batch size is set to 1, meaning that the dataset will be divided into
# ↪batches, each containing 1 example.
# The `collate_fn` function is passed to handle the batching and padding of
# ↪sequences.
tf_train_set = model.prepare_tf_dataset(
    tokenized_train_ubuntu_dialogue_ds,
    shuffle=True,
    batch_size=1,
    collate_fn=data_collator,
)

# Preparing TensorFlow dataset for validation:
# Similar to the training set preparation, but the dataset is not shuffled in
# ↪this case.
# This dataset will be used for validating the model's performance during
# ↪training.
tf_validation_set = model.prepare_tf_dataset(
    tokenized_validation_ubuntu_dialogue_ds,
    shuffle=False,
    batch_size=1,
    collate_fn=data_collator,
)

# Preparing TensorFlow dataset for testing:
# Similar to the validation set preparation, this dataset will be used for
# ↪evaluating the model's performance after training.
tf_test_set = model.prepare_tf_dataset(
    tokenized_test_ubuntu_dialogue_ds,
    shuffle=False,
    batch_size=1,
    collate_fn=data_collator,
)
```

```
[89]: # Define a list of optimizers
optimizers = ['adam', 'sgd', 'adagrad']
```

```
[90]: # Parameters for capturing the best optimizer used in the model
best_eval_loss = float('inf')
best_model = None
best_optimizer = None
```

0.0.8 Model Compile, Train and Evaluate

- Iterate over the optimizers to determine the best one to be used
- compile the model
- train the model on the train and validation datasets
- evaluate the model based on loss leader of optimizer

```
[93]: # Iterate over the optimizers
for optimizer in optimizers:
    if optimizer == 'adam':
        opt = tf.keras.optimizers.Adam(learning_rate=2e-5, weight_decay=0.01) #_
        ↪decay = weight_decay
    elif optimizer == 'adafactor':
        opt = tf.keras.optimizers.Adafactor(learning_rate=2e-5, weight_decay=0.
        ↪01)
    elif optimizer == 'sgd':
        opt = tf.keras.optimizers.SGD(learning_rate=2e-5, weight_decay=0.01)
    elif optimizer == 'adagrad':
        opt = tf.keras.optimizers.Adagrad(learning_rate=2e-5, weight_decay=0.01)

    model.compile(optimizer=opt) # can also add metrics here

    # Train the model
    model.fit(x=tf_train_set, validation_data=tf_validation_set, epochs=3) #_
    ↪callbacks=callbacks

    # Evaluate the model
    eval_loss = model.evaluate(tf_validation_set) # callbacks=callbacks

    # If this model has a lower evaluation loss than the current best, update_
    ↪the best model and optimizer
    if eval_loss < best_eval_loss:
        best_eval_loss = eval_loss
        best_model = model
        best_optimizer = optimizer
```

WARNING:absl:At this time, the v2.11+ optimizer `tf.keras.optimizers.Adam` runs slowly on M1/M2 Macs, please use the legacy Keras optimizer instead, located at `tf.keras.optimizers.legacy.Adam`.

WARNING:absl:There is a known slowdown when using v2.11+ Keras optimizers on M1/M2 Macs. Falling back to the legacy Keras optimizer, i.e.,
`tf.keras.optimizers.legacy.Adam`.

Epoch 1/3

2023-10-15 23:05:50.762160: W

tensorflow/tsl/platform/profile_utils/cpu_utils.cc:128] Failed to get CPU frequency: 0 Hz

10/10 [=====] - 25s 1s/step - loss: 1.1745 - val_loss: 0.4351

Epoch 2/3

10/10 [=====] - 9s 966ms/step - loss: 0.3026 - val_loss: 0.1261

Epoch 3/3

10/10 [=====] - 9s 976ms/step - loss: 0.0799 - val_loss: 0.0415

10/10 [=====] - 3s 294ms/step - loss: 0.0415

WARNING:absl:At this time, the v2.11+ optimizer `tf.keras.optimizers.SGD` runs slowly on M1/M2 Macs, please use the legacy Keras optimizer instead, located at `tf.keras.optimizers.legacy.SGD`.

WARNING:absl:There is a known slowdown when using v2.11+ Keras optimizers on M1/M2 Macs. Falling back to the legacy Keras optimizer, i.e.,
`tf.keras.optimizers.legacy.SGD`.

Epoch 1/3

10/10 [=====] - 22s 2s/step - loss: 0.0454 - val_loss: 0.0413

Epoch 2/3

10/10 [=====] - 10s 1000ms/step - loss: 0.0515 - val_loss: 0.0412

Epoch 3/3

10/10 [=====] - 10s 1s/step - loss: 0.0465 - val_loss: 0.0411

10/10 [=====] - 3s 276ms/step - loss: 0.0411

WARNING:absl:At this time, the v2.11+ optimizer `tf.keras.optimizers.Adagrad` runs slowly on M1/M2 Macs, please use the legacy Keras optimizer instead, located at `tf.keras.optimizers.legacy.Adagrad`.

WARNING:absl:There is a known slowdown when using v2.11+ Keras optimizers on M1/M2 Macs. Falling back to the legacy Keras optimizer, i.e.,
`tf.keras.optimizers.legacy.Adagrad`.

Epoch 1/3

10/10 [=====] - 23s 1s/step - loss: 0.0462 - val_loss: 0.0407

Epoch 2/3

10/10 [=====] - 10s 1s/step - loss: 0.0453 - val_loss: 0.0403

```
Epoch 3/3
10/10 [=====] - 9s 957ms/step - loss: 0.0429 -
val_loss: 0.0399
10/10 [=====] - 2s 239ms/step - loss: 0.0399
```

```
[94]: print('Best Model to use based on the optimizer: ', best_model)
      print('Best Optimizer to use based loss leader comparison: ', best_optimizer)
```

```
Best Model to use based on the optimizer: <transformers.models.blenderbot.model
ing_tf_blenderbot.TFBlenderbotForConditionalGeneration object at 0x5990eb0d0>
Best Optimizer to use based loss leader comparison: adagrad
```

0.0.9 Model Evaluate

- evaluate the model's performance using the model evaluate and model predict

```
[95]: eval_results = best_model.evaluate(tf_validation_set)
      print(f'Perplexity: {math.exp(eval_results):.2f}')
```

```
10/10 [=====] - 2s 233ms/step - loss: 0.0399
Perplexity: 1.04
```

```
[134]: acc = evaluate.load('accuracy')
      f1_score = evaluate.load('f1')
      precis = evaluate.load('precision')
      recalll = evaluate.load('recall')
```

```
[97]: # Round the predictions to turn them into "0" or "1" labels
      test_preds = model.predict(tf_validation_set)
```

```
10/10 [=====] - 6s 232ms/step
```

```
[116]: predicted_ids = tf.argmax(test_preds.logits, axis=-1)
      predicted_ids_tensor = tf.constant(predicted_ids, dtype=tf.int64)
      references = tf.constant(tokenized_validation_ubuntu_dialogue_ds['labels'],
      ↪dtype=tf.int64)

      predicted_ids_numpy = predicted_ids_tensor.numpy().flatten()
      references_numpy = references.numpy().flatten()
```

```
[113]: accuracy_score = acc.compute(predictions=predicted_ids_numpy,
      ↪references=references_numpy)

      print('Accuracy Score: ', accuracy_score)
```

```
Accuracy Score: {'accuracy': 0.9907692307692307}
```

```
[ ]:
```

```
[135]: # Calculating precision score:
# The `compute` method is called on the precision object (`precis`) to
# calculate the precision score.
# Predictions and references are converted to numpy arrays and flattened before
# being passed to the `compute` method.
# The 'weighted' average is used to calculate a single precision score that
# takes class imbalance into account.
precision_score = precis.compute(predictions=predicted_ids_tensor.numpy().
# flatten(), references=references.numpy().flatten(), average='weighted')

# Calculating recall score:
# Similarly, the `compute` method is called on the recall object (`recalll`) to
# calculate the recall score.
# The 'weighted' average is used to calculate a single recall score that takes
# class imbalance into account.
recall_score = recalll.compute(predictions=predicted_ids_tensor.numpy().
# flatten(), references=references.numpy().flatten(), average='weighted')

# Calculating F1 score:
# The `compute` method is called on the F1 score object (`f1_score`) to
# calculate the F1 score.
# The 'weighted' average is used to calculate a single F1 score that takes
# class imbalance into account.
f1_score = f1_score.compute(predictions=predicted_ids_tensor.numpy().flatten(),
# references=references.numpy().flatten(), average='weighted')

# Printing the calculated scores:
print('Precision Score: ', precision_score)
print('Recall Score: ', recall_score)
print('F1 Score: ', f1_score)
```

```
precision Score: {'precision': 0.9865765004226542}
recall Score: {'recall': 0.9907692307692307}
F1 Score: {'f1': 0.9878009561060409}
```

```
/Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-
packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning:
Precision is ill-defined and being set to 0.0 in labels with no predicted
samples. Use `zero_division` parameter to control this behavior.
  _warn_prf(average, modifier, msg_start, len(result))
/Users/hudsonshimanyula/anaconda3/envs/AI_MASTERS_ENV/lib/python3.11/site-
packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning: Recall
is ill-defined and being set to 0.0 in labels with no true samples. Use
`zero_division` parameter to control this behavior.
  _warn_prf(average, modifier, msg_start, len(result))
```


0.0.10 Model Save

```
[136]: # Define the directory where the files are located
save_model_directory = "./bestModel"
```

```
[137]: ### Save the weights of the model
best_model.save_pretrained(save_model_directory)
```

0.0.11 Perform Inference

- recall the saved model
- check the model's architecture
- test the model to confirm it can perform causal inference calculations properly

```
[5]: from transformers import BlenderbotTokenizer, \
    ↪TFBlenderbotForConditionalGeneration
import tensorflow as tf

# Load the model and tokenizer from the directory where you saved your model
save_model_directory = "./bestModel"
tokenizer = BlenderbotTokenizer.from_pretrained('facebook/
    ↪blenderbot-400M-distill')
model= TFBlenderbotForConditionalGeneration.
    ↪from_pretrained(save_model_directory)

def get_response(user_input):
    # Tokenize the user input
    inputs = tokenizer(user_input, return_tensors='tf').input_ids
    # Generate a response
    outputs = model.generate(input_ids=inputs, max_new_tokens=100, \
    ↪do_sample=True, top_k=30, top_p=0.85)
    # Decode the response
    response = tokenizer.batch_decode(outputs, skip_special_tokens=True)[0]
    return response

def print_welcome_message():
    print("=====")
    print("  Welcome to the Chatbot Interface!  ")
    print(" Type 'quit', 'exit', or 'bye' to exit ")
    print("=====")

print_welcome_message()

# Interactive loop
while True:
    user_input = input("You: ")
```

```

if user_input.lower() in ['quit', 'exit', 'bye']:
    print("Bot: Bye! Have a great day.")
    break
print(f"You: {user_input}")
response = get_response(user_input)
print(f"Bot: {response}")

```

All model checkpoint layers were used when initializing
TFBlenderbotForConditionalGeneration.

All the layers of TFBlenderbotForConditionalGeneration were initialized from the
model checkpoint at ./bestModel.

If your task is similar to the task the model of the checkpoint was trained on,
you can already use TFBlenderbotForConditionalGeneration for predictions without
further training.

=====

Welcome to the Chatbot Interface!

Type 'quit', 'exit', or 'bye' to exit

=====

You: Hello

Bot: Hi, how are you doing today? I'm just relaxing before work, how about you?

You: where do you work

Bot: I work at a grocery store. It's not the best but it pays the bills.

You: Well thats good. Straight to the point, I need help with my linux system

Bot: I know how you feel. I need to get a new one as well.

You:

Bot: I'm not sure what you mean by that, but I hope you have a great day!

You: Can you help me troubleshoot an issue

Bot: Sure, what is the issue? I can help you if you want to talk about it.

Bot: Bye! Have a great day.